

```
1 !pip install --upgrade transformers
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (4.47.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers) (3.
Requirement already satisfied: huggingface-hub<1.0,>=0.24.0 in /usr/local/lib/python3.10/dist-packages (fr
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transforme
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transfor
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from transformers) (2.
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/python3.10/dist-packages (from tra
Requirement already satisfied: safetensors>=0.4.1 in /usr/local/lib/python3.10/dist-packages (from transfc
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (from transformers) (
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.10/dist-packages (from huggingfa
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from r
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->tra
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from request
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from request
```

```
1 !pip install numpy==1.23.5
```

```
Requirement already satisfied: numpy==1.23.5 in /usr/local/lib/python3.10/dist-packages (1.23.5)
```

```
1 import torch
```

```
1 import kagglehub
```

```
2
```

```
3 # Download latest version
```

```
4 path = kagglehub.dataset_download("pkdarabi/the-drug-name-detection-dataset")
```

```
5
```

```
6 print("Path to dataset files:", path)
```

```
Path to dataset files: /root/.cache/kagglehub/datasets/pkdarabi/the-drug-name-detection-dataset/versions/1
```

```
1 # Dataset paths
```

```
2 data_dir = "/root/.cache/kagglehub/datasets/pkdarabi/the-drug-name-detection-dataset/versions/1"
```

```
1 from torchvision import datasets, transforms
```

```
2 # Transformations
```

```
3 transform = transforms.Compose([
```

```
4     transforms.Resize((224, 224)),
```

```
5     transforms.ToTensor(),
```

```
6     transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
```

```
7 ])
```

```
1 # Load dataset
```

```
2 dataset = datasets.ImageFolder(data_dir, transform=transform)
```

```
1 # Split dataset
```

```
2 train_size = int(0.7 * len(dataset))
```

```
3 val_size = int(0.2 * len(dataset))
```

```
4 test_size = len(dataset) - train_size - val_size
```

```
5 train_dataset, val_dataset, test_dataset = torch.utils.data.random_split(
```

```

6     dataset, [train_size, val_size, test_size]
7 )

1 from torch.utils.data import DataLoader
2
3 # Create dataloaders
4 train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
5 val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False)
6 test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

1 import matplotlib.pyplot as plt
2 import numpy as np
3
4 # Function to visualize a few images from the dataset
5 def show_sample_images(loader, classes):
6     data_iter = iter(loader)
7     images, labels = next(data_iter)
8     images = images[:6] # Show first 6 images
9     labels = labels[:6]
10
11     fig, axes = plt.subplots(1, 6, figsize=(15, 5))
12     for i in range(6):
13         img = images[i].permute(1, 2, 0).numpy() * 0.5 + 0.5 # Denormalize
14         axes[i].imshow(img)
15         axes[i].set_title(classes[labels[i]])
16         axes[i].axis('off')
17     plt.show()
18
19 # Visualize samples
20 show_sample_images(train_loader, dataset.classes)
21

```



```

1 train_transform = transforms.Compose([
2     transforms.RandomHorizontalFlip(),
3     transforms.RandomRotation(15),
4     transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1),
5     transforms.Resize((224, 224)),
6     transforms.ToTensor(),
7     transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
8 ])
9
10 train_dataset.dataset.transform = train_transform

1 from collections import Counter
2
3 # Get class distribution
4 class_counts = Counter([label for _, label in dataset.samples])
5 for cls, count in class_counts.items():

```

```
6 print(f"{dataset.classes[cls]}: {count}")
7
```

```
⇒ test: 182
   train: 1276
   valid: 365
```

```
1 pip install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cu118
```

```
⇒ Looking in indexes: https://download.pytorch.org/whl/cu118
Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (2.5.1+cu121)
Requirement already satisfied: torchvision in /usr/local/lib/python3.10/dist-packages (0.20.1+cu121)
Requirement already satisfied: torchaudio in /usr/local/lib/python3.10/dist-packages (2.5.1+cu121)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch) (3.16.1)
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages (from torch) (4.12.2)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch) (2.8.8)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch) (3.1.4)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch) (2024.10.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (from torch) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy==1.13.1) (1.3.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from torchvision) (1.23.5)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.10/dist-packages (from torchvision) (10.4.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch) (2.1.5)
```

```
1 import numpy
2 from transformers import CLIPProcessor, CLIPModel
3 import torch
```

```
1 # Step 2 : Fine - Tune the CLIP Model
2 from torch.optim import AdamW
3 from transformers import CLIPProcessor, CLIPModel
4 import torch
5 from torch.optim import Adam
6 from torch.nn import CrossEntropyLoss
7 import torch.nn as nn
8
9 # Load pre-trained CLIP
10 clip_model = CLIPModel.from_pretrained("openai/clip-vit-base-patch32")
11 clip_processor = CLIPProcessor.from_pretrained("openai/clip-vit-base-patch32")
12
13 # for param in clip_model.parameters():
14 #     param.requires_grad = False
15
16 # # Fine-tune only the classification layer or specific layers
17 # for param in clip_model.text_projection.parameters():
18 #     param.requires_grad = True
19
20 # # Enable gradient computation for the vision encoder
21 # for param in clip_model.vision_model.parameters():
22 #     param.requires_grad = True # Fine-tune all layers of the visual encoder
23
24 # # If you also want to fine-tune the text encoder:
25 # for param in clip_model.text_model.parameters():
26 #     param.requires_grad = True # Fine-tune all layers of the text encoder
27
28 # Fine-tuning settings
29 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
30 clip_model.to(device)
```

```

31 optimizer = Adam(clip_model.parameters(), lr=1e-4)
32 criterion = nn.CrossEntropyLoss()

1 # De-normalize images for CLIPProcessor
2 def denormalize(images, mean, std):
3     mean = torch.tensor(mean).view(1, 3, 1, 1).to(images.device)
4     std = torch.tensor(std).view(1, 3, 1, 1).to(images.device)
5     return images * std + mean
6
7 epochs = 10
8
9 # Training loop
10 for epoch in range(epochs):
11     clip_model.train()
12     for images, labels in train_loader:
13         # De-normalize images
14         images = denormalize(images, mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
15
16         # Convert images to PIL-compatible format using CLIPProcessor
17         inputs = clip_processor(images=images, return_tensors="pt", do_rescale=False).to(device)
18         labels = labels.to(device)
19
20         # Forward pass
21         image_features = clip_model.get_image_features(pixel_values=inputs["pixel_values"])
22         loss = criterion(image_features, labels)
23
24         # Backward pass
25         optimizer.zero_grad()
26         loss.backward()
27         optimizer.step()
28
29     print(f"Epoch {epoch + 1}/{epochs} - Loss: {loss.item()}")

```

```

→ Epoch 1/10 - Loss: 0.9677035808563232
Epoch 2/10 - Loss: 1.0617669820785522
Epoch 3/10 - Loss: 0.9015187621116638
Epoch 4/10 - Loss: 0.9433690309524536
Epoch 5/10 - Loss: 0.77410888671875
Epoch 6/10 - Loss: 0.7217981219291687
Epoch 7/10 - Loss: 0.7909273505210876
Epoch 8/10 - Loss: 0.8536691069602966
Epoch 9/10 - Loss: 0.7055936455726624
Epoch 10/10 - Loss: 0.8445102572441101

```

```

1 from sklearn.metrics import accuracy_score
2 import torch
3
4 # Ensure the model is in evaluation mode
5 clip_model.eval()
6 all_preds, all_labels = [], []
7
8 with torch.no_grad():
9     for images, labels in test_loader:
10         # De-normalize images for CLIPProcessor
11         images = denormalize(images, mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
12
13         # Clamp pixel values to [0, 1]
14         images = torch.clamp(images, min=0.0, max=1.0)
15
16         # Process images

```

```

17     inputs = clip_processor(images=images.permute(0, 2, 3, 1).cpu(), return_tensors="pt", do_rescale=False)
18     pixel_values = inputs["pixel_values"].to(device)
19
20     # Extract image features and classify
21     image_features = clip_model.get_image_features(pixel_values=pixel_values)
22
23     # Use a classification layer if added during training
24     preds = torch.argmax(image_features, dim=-1)
25
26     # Store predictions and true labels
27     all_preds.extend(preds.cpu().numpy())
28     all_labels.extend(labels.numpy())
29
30 # Calculate accuracy
31 accuracy = accuracy_score(all_labels, all_preds)
32 print(f"Image Recognition Accuracy: {accuracy * 100:.2f}%")
33
34

```

➞ Image Recognition Accuracy: 72.13%

```

1 from TTS.api import TTS
2
3 # Initialize Tacotron2
4 tts = TTS(model_name="tts_models/en/ljspeech/tacotron2-DDC", gpu=torch.cuda.is_available())
5
6 def generate_voice_output(medicine_name, output_path="medicine_output.wav"):
7     tts.tts_to_file(medicine_name, file_path=output_path)
8     print(f"Voice output saved as {output_path}")
9
10 # Example usage
11 generate_voice_output("Paracetamol")
12

```

➞ /usr/local/lib/python3.10/dist-packages/TTS/api.py:70: UserWarning: `gpu` will be deprecated. Please use `tts.to(device)` instead.

```

> warnings.warn("`gpu` will be deprecated. Please use `tts.to(device)` instead.")
> tts_models/en/ljspeech/tacotron2-DDC is already downloaded.
> vocoder_models/en/ljspeech/hifigan_v2 is already downloaded.
> Using model: Tacotron2
> Setting up Audio Processor...
| > sample_rate:22050
| > resample:False
| > num_mels:80
| > log_func:np.log
| > min_level_db:-100
| > frame_shift_ms:None
| > frame_length_ms:None
| > ref_level_db:20
| > fft_size:1024
| > power:1.5
| > preemphasis:0.0
| > griffin_lim_iters:60
| > signal_norm:False
| > symmetric_norm:True
| > mel_fmin:0
| > mel_fmax:8000.0
| > pitch_fmin:1.0
| > pitch_fmax:640.0
| > spec_gain:1.0
| > stft_pad_mode:reflect
| > max_norm:4.0
| > clip_norm:True
| > do_trim_silence:True

```

```

| > trim_db:60
| > do_sound_norm:False
| > do_amp_to_db_linear:True
| > do_amp_to_db_mel:True
| > do_rms_norm:False
| > db_level:None
| > stats_path:None
| > base:2.718281828459045
| > hop_length:256
| > win_length:1024
/usr/local/lib/python3.10/dist-packages/TTS/utils/io.py:54: FutureWarning: You are using `torch.load` w
return torch.load(f, map_location=map_location, **kwargs)
> Model's reduction rate `r` is set to: 1
> Vocoder Model: hifigan
> Setting up Audio Processor...
| > sample_rate:22050
| > resample:False
| > num_mels:80
| > log_func:np.log
| > min_level_db:-100
| > frame_shift_ms:None
| > frame_length_ms:None
| > ref_level_db:20
| > fft_size:1024
| > power:1.5
| > preemphasis:0.0
| > griffin_lim_iters:60
| > signal_norm:False

```

```

1 import numpy as np
2
3 # Example listener scores
4 listener_scores = {
5     "listener_1": [5, 4, 4],
6     "listener_2": [4, 4, 5],
7     "listener_3": [5, 5, 4]
8 }
9
10 # Calculate Mean Opinion Score
11 mos_score = np.mean([np.mean(scores) for scores in listener_scores.values()])
12 print(f"Mean Opinion Score (MOS): {mos_score:.1f}")

```

➡ Mean Opinion Score (MOS): 4.4

```
1 pip install gTTS
```

➡ Requirement already satisfied: gTTS in /usr/local/lib/python3.10/dist-packages (2.5.4)  
Requirement already satisfied: requests<3,>=2.27 in /usr/local/lib/python3.10/dist-packages (from gTTS) (2.31.0)  
Requirement already satisfied: click<8.2,>=7.1 in /usr/local/lib/python3.10/dist-packages (from gTTS) (8.1.7)  
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.27) (3.3.2)  
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.27) (3.6)  
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.27) (2.2.3)  
Requirement already satisfied: certifi<2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.27) (2024.7.4)

```
1 !pip install TTS
```

➡ Requirement already satisfied: TTS in /usr/local/lib/python3.10/dist-packages (0.22.0)  
Requirement already satisfied: cython>=0.29.30 in /usr/local/lib/python3.10/dist-packages (from TTS) (3.0.12)  
Requirement already satisfied: scipy>=1.11.2 in /usr/local/lib/python3.10/dist-packages (from TTS) (1.13.1)  
Requirement already satisfied: torch>=2.1 in /usr/local/lib/python3.10/dist-packages (from TTS) (2.5.1+cu121)  
Requirement already satisfied: torchaudio in /usr/local/lib/python3.10/dist-packages (from TTS) (2.5.1+cu121)

```

Requirement already satisfied: soundfile>=0.12.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (0.12.0)
Requirement already satisfied: librosa>=0.10.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (0.10.0)
Requirement already satisfied: scikit-learn>=1.3.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (1.3.0)
Requirement already satisfied: inflect>=5.6.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (7.0.0)
Requirement already satisfied: tqdm>=4.64.1 in /usr/local/lib/python3.10/dist-packages (from TTS) (4.66.0)
Requirement already satisfied: anyascii>=0.3.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (0.3.0)
Requirement already satisfied: pyyaml>=6.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (6.0.2)
Requirement already satisfied: fsspec>=2023.6.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (2023.6.0)
Requirement already satisfied: aiohttp>=3.8.1 in /usr/local/lib/python3.10/dist-packages (from TTS) (3.8.1)
Requirement already satisfied: packaging>=23.1 in /usr/local/lib/python3.10/dist-packages (from TTS) (23.1)
Requirement already satisfied: flask>=2.0.1 in /usr/local/lib/python3.10/dist-packages (from TTS) (3.0.0)
Requirement already satisfied: pysbd>=0.3.4 in /usr/local/lib/python3.10/dist-packages (from TTS) (0.3.4)
Requirement already satisfied: umap-learn>=0.5.1 in /usr/local/lib/python3.10/dist-packages (from TTS) (0.5.1)
Requirement already satisfied: pandas<2.0,>=1.4 in /usr/local/lib/python3.10/dist-packages (from TTS) (1.5.3)
Requirement already satisfied: matplotlib>=3.7.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (3.7.0)
Requirement already satisfied: trainer>=0.0.32 in /usr/local/lib/python3.10/dist-packages (from TTS) (0.0.32)
Requirement already satisfied: coqpit>=0.0.16 in /usr/local/lib/python3.10/dist-packages (from TTS) (0.0.16)
Requirement already satisfied: jieba in /usr/local/lib/python3.10/dist-packages (from TTS) (0.42.1)
Requirement already satisfied: pypinyin in /usr/local/lib/python3.10/dist-packages (from TTS) (0.53.0)
Requirement already satisfied: hangul-romanize in /usr/local/lib/python3.10/dist-packages (from TTS) (0.0.1)
Requirement already satisfied: gruut==2.2.3 in /usr/local/lib/python3.10/dist-packages (from gruut[de,en,es,fr,ja,pt,ru]) (2.2.3)
Requirement already satisfied: jamo in /usr/local/lib/python3.10/dist-packages (from TTS) (0.4.1)
Requirement already satisfied: nltk in /usr/local/lib/python3.10/dist-packages (from TTS) (3.9.1)
Requirement already satisfied: g2pkk>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from TTS) (0.1.1)
Requirement already satisfied: bangla in /usr/local/lib/python3.10/dist-packages (from TTS) (0.0.2)
Requirement already satisfied: bnnumerizer in /usr/local/lib/python3.10/dist-packages (from TTS) (0.0.2)
Requirement already satisfied: bnunicodenormalizer in /usr/local/lib/python3.10/dist-packages (from TTS) (0.0.2)
Requirement already satisfied: einops>=0.6.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (0.8.0)
Requirement already satisfied: transformers>=4.33.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (4.33.0)
Requirement already satisfied: encodec>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from TTS) (0.1.1)
Requirement already satisfied: unidecode>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from TTS) (1.3.2)
Requirement already satisfied: num2words in /usr/local/lib/python3.10/dist-packages (from TTS) (0.5.13)
Requirement already satisfied: spacy>=3 in /usr/local/lib/python3.10/dist-packages (from spacy[ja]>=3->3.7.4) (3.7.4)
Requirement already satisfied: numpy==1.22.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (1.22.0)
Requirement already satisfied: numba>=0.57.0 in /usr/local/lib/python3.10/dist-packages (from TTS) (0.57.0)
Requirement already satisfied: Babel<3.0.0,>=2.8.0 in /usr/local/lib/python3.10/dist-packages (from gruut[de,en,es,fr,ja,pt,ru]) (2.10.1)
Requirement already satisfied: dateparser~=1.1.0 in /usr/local/lib/python3.10/dist-packages (from gruut[de,en,es,fr,ja,pt,ru]) (1.1.8)
Requirement already satisfied: gruut-ipa<1.0,>=0.12.0 in /usr/local/lib/python3.10/dist-packages (from gruut[de,en,es,fr,ja,pt,ru]) (0.12.0)
Requirement already satisfied: gruut-lang-en~=2.0.0 in /usr/local/lib/python3.10/dist-packages (from gruut[de,en,es,fr,ja,pt,ru]) (2.0.0)
Requirement already satisfied: jsonlines~=1.2.0 in /usr/local/lib/python3.10/dist-packages (from gruut[de,en,es,fr,ja,pt,ru]) (1.2.0)
Requirement already satisfied: networkx<3.0.0,>=2.5.0 in /usr/local/lib/python3.10/dist-packages (from gruut[de,en,es,fr,ja,pt,ru]) (2.8.8)
Requirement already satisfied: python-crfsuite~=0.9.7 in /usr/local/lib/python3.10/dist-packages (from gruut[de,en,es,fr,ja,pt,ru]) (0.9.7)
Requirement already satisfied: gruut-lang-de~=2.0.0 in /usr/local/lib/python3.10/dist-packages (from gruut[de,en,es,fr,ja,pt,ru]) (2.0.0)
Requirement already satisfied: gruut-lang-fr~=2.0.0 in /usr/local/lib/python3.10/dist-packages (from gruut[de,en,es,fr,ja,pt,ru]) (2.0.0)
Requirement already satisfied: gruut-lang-es~=2.0.0 in /usr/local/lib/python3.10/dist-packages (from gruut[de,en,es,fr,ja,pt,ru]) (2.0.0)
Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (2.3.0)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (1.1.2)
Requirement already satisfied: async-timeout<6.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (4.0.3)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (23.1.0)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (1.1.1)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (6.0.0)
Requirement already satisfied: yarl>=1.8.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (1.8.0)

```

```

1 # Resnet
2
3 import torch
4 import torch.nn as nn
5 from torchvision import models, transforms
6 from torch.utils.data import DataLoader
7
8 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
9
10 # Load ResNet-50
11 class ResNetClassifier(nn.Module):
12     def __init__(self, num_classes):

```

```

13     super(ResNetClassifier, self).__init__()
14     self.backbone = models.resnet50(pretrained=True)
15     self.backbone.fc = nn.Linear(self.backbone.fc.in_features, num_classes)
16
17     def forward(self, x):
18         return self.backbone(x)
19
20 # Initialize ResNet model
21 # Use the dataset object to get the number of classes
22 resnet_model = ResNetClassifier(num_classes=len(dataset.classes)).to(device)
23 resnet_optimizer = torch.optim.Adam(resnet_model.parameters(), lr=1e-4)
24 resnet_criterion = nn.CrossEntropyLoss()

```

```

1 # Training loop for ResNet
2 epochs = 10
3 for epoch in range(epochs):
4     resnet_model.train()
5     for images, labels in train_loader:
6         images, labels = images.to(device), labels.to(device)
7
8         # Forward pass
9         outputs = resnet_model(images)
10        loss = resnet_criterion(outputs, labels)
11
12        # Backward pass
13        resnet_optimizer.zero_grad()
14        loss.backward()
15        resnet_optimizer.step()
16
17    print(f"Epoch {epoch+1}/{epochs} - Loss: {loss.item()}")

```

```

→ Epoch 1/10 - Loss: 0.5917323231697083
Epoch 2/10 - Loss: 0.7319120168685913
Epoch 3/10 - Loss: 0.5736165642738342
Epoch 4/10 - Loss: 0.5468763709068298
Epoch 5/10 - Loss: 0.3438940942287445
Epoch 6/10 - Loss: 0.23517988622188568
Epoch 7/10 - Loss: 0.3595772385597229
Epoch 8/10 - Loss: 0.1950729489326477
Epoch 9/10 - Loss: 0.19309721887111664
Epoch 10/10 - Loss: 0.19073885679244995

```

```

1 from sklearn.metrics import accuracy_score
2
3 resnet_model.eval()
4 all_preds, all_labels = [], []
5
6 with torch.no_grad():
7     for images, labels in test_loader:
8         images, labels = images.to(device), labels.to(device)
9         outputs = resnet_model(images)
10        preds = torch.argmax(outputs, dim=-1)
11
12        all_preds.extend(preds.cpu().numpy())
13        all_labels.extend(labels.cpu().numpy())
14
15 accuracy_resnet = accuracy_score(all_labels, all_preds)
16 print(f"ResNet Image Recognition Accuracy: {accuracy_resnet * 100:.2f}%")
17

```



➞ ResNet Image Recognition Accuracy: 56.28%

```
1 from gtts import gTTS
2 import os
3
4 def basic_tts(medicine_name, output_path="basic_tts_output.wav"):
5     tts = gTTS(text=medicine_name, lang='en')
6     tts.save(output_path)
7     print(f"Basic TTS Voice Output Saved: {output_path}")
8
9 # Example usage
10 basic_tts("Paracetamol")
```

➞ Basic TTS Voice Output Saved: basic\_tts\_output.wav

```
1 import numpy as np
2
3 # Example listener scores
4 listener_scores_resnet = {
5     "listener_1": [3, 4, 3],
6     "listener_2": [4, 4, 3],
7     "listener_3": [3, 4, 4]
8 }
9
10 # Calculate MOS
11 mos_resnet = np.mean([np.mean(scores) for scores in listener_scores_resnet.values()])
12 print(f"Mean Opinion Score (MOS) for ResNet + Basic TTS: {mos_resnet:.1f}")
13
```

➞ Mean Opinion Score (MOS) for ResNet + Basic TTS: 3.6

```
1 # Baseline SOTA
2 !pip install timm
```

➞ Requirement already satisfied: timm in /usr/local/lib/python3.10/dist-packages (1.0.12)  
 Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (from timm) (2.5.1+cu121)  
 Requirement already satisfied: torchvision in /usr/local/lib/python3.10/dist-packages (from timm) (0.20.1+cu121)  
 Requirement already satisfied: pyyaml in /usr/local/lib/python3.10/dist-packages (from timm) (6.0.2)  
 Requirement already satisfied: huggingface\_hub in /usr/local/lib/python3.10/dist-packages (from timm) (0.20.3)  
 Requirement already satisfied: safetensors in /usr/local/lib/python3.10/dist-packages (from timm) (0.4.5)  
 Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from huggingface\_hub->timm) (3.13.1)  
 Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.10/dist-packages (from huggingface\_hub->timm) (2024.6.1)  
 Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.10/dist-packages (from huggingface\_hub->timm) (24.0)  
 Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from huggingface\_hub->timm) (2.32.0)  
 Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.10/dist-packages (from huggingface\_hub->timm) (4.67.1)  
 Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from torch->timm) (4.11.0)  
 Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch->timm) (3.3)  
 Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (from torch->timm) (3.1.4)  
 Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (from torch->timm) (1.13.1)  
 Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy==1.13.1->torch->timm) (1.3.0)  
 Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from torchvision->timm) (1.26.4)  
 Requirement already satisfied: pillow!=8.3.\*,>=5.3.0 in /usr/local/lib/python3.10/dist-packages (from torchvision->timm) (10.4.0)  
 Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2->torch->timm) (2.1.5)  
 Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->torch->timm) (3.3.2)  
 Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->torch->timm) (3.10.1)  
 Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->torch->timm) (2.2.3)  
 Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->torch->timm) (2024.7.4)

```
1 !pip install numpy==1.23.5
```

➞ Requirement already satisfied: numpy==1.23.5 in /usr/local/lib/python3.10/dist-packages (1.23.5)

```
1 import torch.nn as nn
2 from torch.optim import Adam
3 from torchvision import transforms, datasets
4 from transformers import ViTForImageClassification, ViTImageProcessor
5
6 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
7
8 # Initialize Vision Transformer (ViT)
9 model_name = "google/vit-base-patch16-224-in21k"
10 feature_extractor = ViTImageProcessor.from_pretrained(model_name)
11 vit_model = ViTForImageClassification.from_pretrained(model_name, num_labels=len(dataset.classes)).to(device)
12
13 # Loss and Optimizer
14 criterion = nn.CrossEntropyLoss()
15 optimizer = Adam(vit_model.parameters(), lr=1e-4)
16
17 # Training Loop
18 epochs = 10
19 for epoch in range(epochs):
20     vit_model.train()
21     running_loss = 0.0
22     for images, labels in train_loader:
23         # Move images and labels to device
24         images, labels = images.to(device), labels.to(device)
25
26         # Ensure pixel values are in the correct range
27         images = torch.clamp(images, 0, 1)
28
29         # Process images with the feature extractor
30         inputs = feature_extractor(images=images.permute(0, 2, 3, 1).cpu(), return_tensors="pt", do_rescale=False)
31         optimizer.zero_grad()
32
33         # Forward pass
34         outputs = vit_model(**inputs)
35         loss = criterion(outputs.logits, labels)
36
37         # Backward pass
38         loss.backward()
39         optimizer.step()
40         running_loss += loss.item()
41
42     print(f"Epoch {epoch + 1}/{epochs} - Loss: {running_loss/len(train_loader):.4f}")
43
44
```

➞ Some weights of ViTForImageClassification were not initialized from the model checkpoint at google/vit-base-patch16-224-in21k. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
Epoch 1/10 - Loss: 0.8451
Epoch 2/10 - Loss: 0.7899
Epoch 3/10 - Loss: 0.7627
Epoch 4/10 - Loss: 0.6689
Epoch 5/10 - Loss: 0.5378
Epoch 6/10 - Loss: 0.3397
Epoch 7/10 - Loss: 0.1765
Epoch 8/10 - Loss: 0.1051
Epoch 9/10 - Loss: 0.0677
Epoch 10/10 - Loss: 0.0448
```

```

1 from sklearn.metrics import accuracy_score
2
3 vit_model.eval()
4 all_preds, all_labels = [], []
5
6 with torch.no_grad():
7     for images, labels in test_loader:
8         # Move images and labels to the appropriate device
9         images, labels = images.to(device), labels.to(device)
10
11         # Rescale images to the expected range [0, 1]
12         images = torch.clamp((images + 1) / 2, 0, 1) # Convert from [-1, 1] to [0, 1]
13
14         # Convert images for the ViT feature extractor
15         inputs = feature_extractor(images=images.permute(0, 2, 3, 1).cpu(), return_tensors="pt").to(device)
16
17         # Pass through the ViT model
18         outputs = vit_model(**inputs)
19         preds = torch.argmax(outputs.logits, dim=-1)
20
21         # Store predictions and labels
22         all_preds.extend(preds.cpu().numpy())
23         all_labels.extend(labels.cpu().numpy())
24
25 # Calculate accuracy
26 accuracy = accuracy_score(all_labels, all_preds)
27 print(f"ViT Image Recognition Accuracy: {accuracy * 100:.2f}%")
28

```

 ViT Image Recognition Accuracy: 67.76%

+ Code

+ Text

```

1 # Generate Voice Output with gTTS
2 def generate_voice_output_gtts(text, output_path="output.mp3"):
3     tts = gTTS(text=text, lang="en")
4     tts.save(output_path)
5     print(f"Voice output saved at {output_path}")
6
7 # Example Usage
8 medicine_name = "Paracetamol"
9 generate_voice_output_gtts(f"The medicine name is {medicine_name}.", output_path="medicine_name.mp3")
10

```

 Voice output saved at medicine\_name.mp3

```


1 listener_scores = {
2     "listener_1": [4, 5, 4, 4],
3     "listener_2": [5, 4, 4, 5],
4     "listener_3": [4, 4, 5, 4]
5 }
6
7 import numpy as np
8
9 # Calculate MOS score
10 mos_score = np.mean([np.mean(scores) for scores in listener_scores.values()])
11 print(f"Mean Opinion Score (MOS) for ViT: {mos_score:.1f}")
12

```


➞ Mean Opinion Score (MOS) for ViT: 4.3

```
1 import matplotlib.pyplot as plt
2
3 models = ["CLIP + Tacotron2", "ResNet + Basic TTS", "Baseline (SOTA)"]
4 accuracy = [72.13, 56.28, 67.76] # Replace with your accuracy values
5
6 plt.figure(figsize=(8, 6))
7 plt.bar(models, accuracy, color=['blue', 'orange', 'green'])
8 plt.ylabel("Image Recognition Accuracy (%)")
9 plt.title("Accuracy Comparison Across Models")
10 plt.ylim(0, 100)
11 plt.grid(axis="y", linestyle="--", alpha=0.7)
12 plt.show()
13
```

➞



```
1 mos_scores = [4.4, 3.8, 4.3] # Replace with your MOS values
2
3 plt.figure(figsize=(8, 6))
4 plt.bar(models, mos_scores, color=['blue', 'orange', 'green'])
5 plt.ylabel("Mean Opinion Score (MOS)")
6 plt.title("MOS Score Comparison Across Models")
7 plt.ylim(0, 5)
8 plt.grid(axis="y", linestyle="--", alpha=0.7)
9 plt.show()
10
```



1 Start coding or [generate](#) with AI.